Hands-on non-technical tutorial for Bayesian mixed effects regression

Michael Franke & Timo Roettger

Generalized linear mixed models are very versatile and handy tools for statistical inference. Bayesian approaches to applying these models have recently become increasingly popular. This tutorial provides an accessible, non-technical introduction to the use and feel of Bayesian mixed effects regression models. The focus is on data from a factorial-design experiment.

This tutorial should take you about 1 hour.

Motivation & intended audience

This tutorial provides a very basic introduction to Bayesian regression modeling using R (R Core Team, 2017). We wrote this tutorial with a particular reader in mind. If you have used R before and if you have a basic understanding of linear regression and now you want to find out what a Bayesian approach has to offer, this tutorial is for you. In comparison to other introductions (e.g. Sorensen, Hohensteinb, and Vasishth, 2016), this tutorial remains very conceptual. We don't want to "sell Bayes" to you, and we do not want to scare you away with mathematical details. We just want to give you an impression of how a Bayesian regression analysis looks and feels. So no reason to be afraid! But also: no reason to be bored, because we *will* cover all the essential concepts and we *will* explain how to run and interpret the output of a Bayesian regression analysis using the wonderful R package `brms` written by Paul Buerkner (2016).

If you don't have any experience with regression modeling, you will probably still be able to follow but you might also want to consider doing a crash course. To bring you up to speed, we recommend the excellent two-part tutorial by Bodo Winter (2013) on mixed effects regression in a non-Bayesian —a.k.a. classical or frequentist— paradigm. In a sense, this tutorial could be considered part three of Bodo's nice and lofty introduction. We will, for example use the same data set.

This tutorial contains text boxes (with a gray background) which contain additional background information on some topics. The information is sometimes a bit technical but never absolutely necessary for understanding the main ideas. So, feel free to read or skip any of the text boxes to suit your needs.

To follow this tutorial, you should have R installed on your computer (`https://www.r-project.org`). Unless you already have a favorite editor for tinkering with R scripts, we recommend to try out RStudio (`https://www.rstudio.com`). You will also need some packages,

Remember that you can install a package called `XYZ` with the command `install.packages('XYZ')`.

which you can import with the following code:

```
#####################################################
## package includes and options
#####################################################

# package for convenience functions (e.g. plotting)
library(tidyverse)

# package for Bayesian regression modeling
library(brms)
# option for Bayesian regression models: use all available cores
    for parallel computing
options(mc.cores = parallel::detectCores())
```

If you do not want to copy-paste, all code and data for this tutorial is also available for download here: `https://github.com/michael-franke/bayes_mixed_regression_tutorial`

Data, research questions & hypotheses

Imagine we are experimental researchers. Therefore, we collect data to answer questions of interest about how nature works. For example, we might want to know whether voice pitch differs across female and male speakers, and whether it differs across social contexts (say: informal and polite contexts). — To answer our questions, we come up with a nifty experimental design; we lure a group of people into the lab; we ask them to say different words in different social contexts; we record their voices; and we extract numbers from these recordings, for example, the pitch values of their voices. We then want to find out whether our data provide evidence for any assumed relationships. So far so good.

In this tutorial, we are looking at exactly these kind of data (following Winter, 2013). To load the data into your R environment, run the following code:

The data is originally from research presented by Winter and Grawunder (2012)

If you are familiar with the previous tutorials by Winter (2013), note that we 'massaged' the data a bit, i.e., we renamed variables and removed a line with missing data.

```
  # load the data into variable 'politedata'
 politedata = read_csv("https://raw.githubusercontent.com/
    michael-franke/bayes_mixed_regression_tutorial/master/code/
    politeness_data.csv")
```

Type `head(politedata)` and you should see the first lines of the imported data:

Here, we show only part of the output that you can see when executing this command.

```
> head(politedata)
   subject gender sentence context pitch
   <chr>   <chr>  <chr>    <chr>   <dbl>
 1 F1      F      S1       pol      213.
 2 F1      F      S1       inf      204.
 3 F1      F      S2       pol      285.
 4 F1      F      S2       inf      260.
 5 F1      F      S3       pol      204.
```

This data set contains information about different subjects, with an anony-

mous identifier stored in variable `subject`. Because voice pitch is highly dependent on gender (i.e. there are anatomical differences between women and men that affect voice pitch), we stored whether our subjects are F(emale) or M(ale) in variable `gender`. Subjects produced different sentences (stored in variable `sentence`), and the experiment manipulated whether the sentence was produced in a polite or an informal context, indicated by the variable `context`. Crucially, each row contains a measurement of pitch in Hz stored in variable `pitch`.

Often, we are interested in comparing a **dependent variable** (here `pitch`) across different conditions or groups, i.e. **independent variables** (here `gender` and `context`). Before our data collection, we might have formulated concrete predictions about the relationship between the dependent variable and the independent variables. For example, we might have formulated the following three hypotheses: [tr: up till here, this reads really lofty and flows nicely]

H1: Female speakers have a lower average pitch in polite than in informal contexts.

H2: Male speakers have a lower average pitch in polite than in informal contexts.

H3: Male speakers have a lower average pitch in informal than female speakers have in polite contexts.

Notice that our hypotheses are formulated explicitly as comparisons of means / averages. The statistical model we will use indeed compares means, and this is common practice, albeit a particular assumption worth highlighting. Commonly, this assumption is implicit, with e.g. H1 being simply formulated as something like "Female speakers' pitch is lower in polite than informal contexts."

Exploring the data visually

To get a first idea of possible relationships in our data, let's plot them. Figure 1 displays the mean pitch values for each sentence (semi-transparent points) across gender and context. The solid points indicate the average pitch values across all subjects. Looking at the plot, we can see that pitch values from female speakers are generally higher than those from male speakers (points in left column are higher than in the right column). We also see that pitch values in the informal context are slightly higher than those in the polite context (blue points are slightly higher than orange points).

Based on keen eye-balling, we might want to shout: "Eureka! The data confirm all of our hypotheses!" But, of course, we need to be more careful. As Bayesians, we would like to translate the data into an expression of **evidence**: does the data provide evidence for our research hypotheses? Or are the observable differences to meager? – Also, notice that there is quite a lot of variability between different sentences (the semi-transparent dots in Figure 1). For example, some values from the informal condition for female speakers (blue points in left column), are lower than their corresponding polite counterparts. Similarly, there could be quite some differences between individual speakers. Consequently, what we want is precise estimates of potential

Extensive plotting is always recommended to start data analysis. You need to know your data inside out. Pictures often reveal complex relationships much better than numbers can.

The code needed to generated the picture in Figure 1 is not reproduced here, but included in the script in the resources for this tutorial: `https://github.com/michael-franke/bayes_mixed_regression_tutorial`

Figure 1: Basic plot of the data display-ing overall averages (thick points) and averages for individual sentences (smaller semitransparent points)

differences between conditions, alongside a measure of how confident we can be in these estimates.

A regression model for our data

Another way of looking at the data in connection with our research hypothe-ses is displayed in Figure 2. Each cell represents one unique combination of the gender and the context factor, and the table shows the mean pitch value for each cell. [tr: font in the side note is larger than the other notes and it might be positioned wrongly.] Our hypotheses can be related to the compar-ison between some of these cell-based means. H1 makes a statement about the comparison between cells 1 and 2 (the context effect for female speakers); H2 makes a statement about cells 3 and 4 (the context effect for male speak-ers); and H3 makes a statement about cells 2 and 3 (the difference between informal male speakers and polite female speakers).

One way of testing our hypotheses using a Bayesian approach to data analysis, is to ask whether the relevant differences between cell means are *credibly different from zero*. This is jargon for asking whether, given the observed data, we should believe that the relevant cell means are different from each other. The fact that we are now in Bayesian land should hit us when we hear the phrases "credibly different" and the "should believe". The Bayesian approach is about updating beliefs based on observations. And we express our beliefs as as probability distributions. At first, this may appear scary or technically involved. But at the end of the day, the intuitions captured by this approach are arguably very natural, and perhaps easier to understand than the reasoning underlying other approaches to statistical

In technical terms, this table is the **design matrix** of our experiment. We have two factors of interest `context` and `gender`, each with two levels. The table shows each combination of levels of all relevant factors. The cells in this table are therefore also called **design cells**.

Figure 2: Means of each design cell, together with research hypotheses as statements about ordinal relations between cell means.

inference. Let's walk through our Bayesian approach step by step.

First, let us look at the **regression model** we want to use. Our regression models assumes that pitch values observed in each cell are sampled from a population that is normally distribution, where each cell c_i has its own mean μ_i. We are interested in the probability of one cell mean being larger than another cell mean, i.e., the probability that $\mu_i > \mu_j$. Put differently, we are interested in the probability that the difference between μ_i and μ_j is larger than zero: $\mu_i - \mu_j > 0$. Figure 3 illustrates the encoding scheme of our cell means in terms of a regression analysis. It assumes that there is a **reference level** for each factor. Here it is the level `female` for the factor `gender` and the level `informal` for the factor `context`.

All cell means can then be expressed in terms of differences between the **intercept** β_0 which is the cell mean of the reference level (here, the cell mean of female subjects in informal contexts) and deviations from this **reference cell** for each individual factor (β_{male}, and β_{polite}), and a so-called **interaction term** $\beta_{\text{pol\&male}}$. In other words, our regression estimates the mean of the reference level and estimates how much we need to adjust this mean when we change either the context level (C2), the gender level (C3), or both (C4). Every value that the model estimates is called a parameter.

A Bayesian analysis of a (fixed effects) regression model

Based on our model of how the data was generated [tr: I dont think this is clear to the reader], a Bayesian analysis asks: what should we believe about the values of the coefficients $\beta_0, \beta_{\text{pol}}, \beta_{\text{male}}$ and $\beta_{\text{pol\&male}}$?; what values for

While other approaches to data analysis might stress the application of statistical tests, a Bayesian approach puts more emphasis on the fact that all statistical inference resolves around a statistical model; which is usually an assumption about how the data was generated. This model is, almost certainly, always false. But even a false model can be useful; conclusions based on a false model can be meaningful and insightful.[tr: I find this remark confusing and maybe a little too much on the "selling Bayes" side of things? consider deleting]

This is so-called **dummy coding** of the regression coefficients. Other coding schemes exist, but are not discussed here.

Figure 3: Coefficients of a dummy-coded regression model for the factorial 2×2 design.

these parameters are likely, given the data, the assumed model, and our initial beliefs about the parameters?

The R package `brms` (Buerkner, 2016) makes it easy to run Bayesian regression models. It uses a very similar formula syntax as related packages for regression analysis. In our case, we want to regress the dependent variable `pitch` against the independent variables `gender` and `context` and their two-way interaction. This model is expressed by the formula:

```
# formula for (fixed effects) regression model
formulaFE = pitch ~ gender * context
```

The Bayesian model can then be fitted with the function `brm` from the `brms` package. We only need to specify the formula and supply the data:

```
# run regression model in brms
modelFE = brm(
  formula = formulaFE,
  data = politedata
)
```

The `brms` packages uses the probabilistic programming language `Stan` in the background. What `brms` does is translating your syntax into Stan code and executes it. The Stan code is then translated to C++ (hence the message about "compiling C++" when you run this code). Conceptually, Stan obtains samples from the posterior distribution, based on an algorithm called *Hamiltonian Monte Carlo*. This is an instance of a more general class of algorithms, called *Markov Chain Monte Carlo* methods. The purpose of these

These initial beliefs, also called **prior beliefs**, are important to get a Bayesian analysis off the ground; a circumstance which is discussed controversially. For many practical purposes, however, the precise choice of prior is not decisive and tools like the `brms` package which we will use here will default to generically reasonable choices of priors for your model (more on this below).[tr: should we maybe use a consistent typography such that the margin notes always pick up on a bold word in the text?]

Info Box 1 provides some background on prior beliefs, likelihood function and posterior beliefs.

methods is to return representative samples from the posterior distribution (without actually having to compute the exact distribution, which might be intractable) [tr: I suggest deleting the brackets. Difficult to understand and not relevant].

You can type in `modelFE` in order to get a summary of the model fit. It should look much like the following output (note that your values might be slightly different due to stochastic variation in the sampling procedure).

```
> modelFE
 Family: gaussian
  Links: mu = identity; sigma = identity
Formula: pitch ~ gender * context
   Data: politedata (Number of observations: 83)
Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
         total post-warmup samples = 4000

Population-Level Effects:
                 Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
Intercept          260.68      8.07   244.99   276.78       2409 1.00
genderM           -116.09     11.44  -138.37   -93.80       2094 1.00
contextpol         -27.38     11.33   -50.01    -5.98       2092 1.00
genderM:contextpol  15.74     16.38   -17.03    49.13       1831 1.00

Family Specific Parameters:
      Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
sigma    36.15      2.87    30.93    42.40       3652 1.00

Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
is a crude measure of effective sample size, and Rhat is the potential
scale reduction factor on split chains (at convergence, Rhat = 1).
```

This summary looks very much like regression model summaries of non-Bayesian model. Lines 2–5 give us information about the model and the data used. Lines 6 and 7 tell us about the sampling procedure. We have a total of 4000 samples from the posterior distribution, obtained from 4 chains all of which had 2000 iterations but discarded the first 1000 as warmup (more on this below) [tr: chains, iterations, warm-up need to be elaborated]. Lines 9–14 contain information about our parameters of interest, i.e. what we are most interested in for evaluating our hypotheses, We will discuss them in detail below. Lines 16–18 contain information that look similar to those in lines 9–14. This is the estimation of the standard deviation `sigma`, describing the variance of the assumed normal distributions (which describe the distribution of measures in each design cell) [tr: maybe wanna call this the residual error instead?]. Finally, lines 20–22 contain general information about the model fit and the information presented in this summary.

Let us look at lines 9–14 in detail now. What these lines give us is a table with four rows, each of which corresponds to a parameter in the model, namely the coefficients shown in Figure 3. The variable `Intercept` refers to our β_0, which represents the mean of our reference level in cell 1 (female

If the model failed to converge or other problems occurred, you would see an informative message in the last part of this summary.

speakers in polite contexts). The variable `genderM` corresponds to our β_{male}, `contextpol` corresponds to our β_{pol}, and `genderM:contextpol` is the interaction term $\beta_{\text{pol\&male}}$. For each of these parameters, the table contains very useful summary statistics based on the samples returned from the model fit. Here we are interested in the information in columns `Estimate`, that is the estimate of each parameter's mean. We are also interested in `l-95% CI` and `u-95% CI`, which give us the lower and upper bound of the **95% credible interval** for each parameter. More details about the information given in the other columns can be found in Info Box 2.

For our purposes, the information about 95% credible intervals is most interesting. Take the parameter `contextpol`, corresponding to our coefficient β_{pol}. This parameter corresponds to the estimated adjustment of the mean of the reference level when we change the context level to polite. The 95% CI is roughly [-50;-6]. This means that we would take values outside of this interval to be sufficiently unlikely to consider plausible values. This means that a very special parameter value for this parameter is implausible, namely 0. In other words, this analysis suggests that we should no believe that 0 is a very likely value for the coefficient β_{pol}; rather we should believe that β_{pol} is most likely negative (based on the priors, the regression model, and the data; all of which may be critically re-assessed if there are good reasons for it[tr: again, this nuance might be too much for the basic tutorial or am I too fluffy?]). — Hurray! This directly addresses our first research hypothesis. In a research paper we could now write: "Based on the regression model, the data suggests that H1 is likely true."

How likely is it that β_{pol} is smaller than 0? —Instead of simply making a binary thumbs-up / thumbs-down decision, it would be even cooler, if we could put a number to it. As Bayesians, we fortunately can. To see how this works, let us have a more intimate look at the samples that the `brm` function returns. We can access the samples of a model fitted with `brm` with the function `posterior_samples`:

```
# extract posterior samples
post_samples_FE = posterior_samples(modelFE)
head(post_samples_FE)
```

The output of this could look like this:

```
> head(post_samples_FE)
  b_Intercept b_genderM b_contextpol b_genderM:contextpol   sigma     lp__
1     255.3 -106.4     -24.7                 15.9 34.3 -420.1
2     252.5 -118.7     -15.7                 22.7 35.4 -421.2
3     254.1 -119.8     -15.9                 15.5 35.7 -420.6
4     270.1 -114.0     -33.7                 26.9 36.3 -423.1
5     275.4 -122.9     -37.1                 23.5 37.5 -422.0
6     281.4 -135.4     -43.4                 25.5 38.9 -423.3
```

Intuitively, the 95% credible interval is the range of values that we can often practically consider credible enough to care about. Formally, the 95% credible interval is the set of convex intervals of parameter values such that the probability density over all intervals sums to 0.95 while no parameter value not included in any interval has higher probability density than any point within.[tr: consider deleting the last sentence. Not relevant for this tutorial, Let's stick to practical aspects]

What you see here is the top 6 rows of a data frame with columns for each parameter and 4000 rows, corresponding to each sample of that parameter (so our sampling method has generated 4000 likely values for each parameter). We can use these samples to produce a density plot. The plot in Figure 4 shows, for each of the four main model parameters an estimate of the posterior density. This is, intuitively put, a plot of how much (relative) credence we, as rational agents, should assign to any particular value of each parameter, given the regression model, the specified prior beliefs, and the data [tr: and now in a way that your daughter understands it ;)]. We see that our beliefs concerning plausible values for the mean of cell 1 (female speakers in polite contexts, the reference cell) should hover around 261. We also see that all values that receive substantial probability density for `contextpol` (our β_{pol}) are negative (as captured in the 95% CI discussed above). Zero is estimated to be a rather unlikely value for this parameter.

The column `lp__` contains the log-probability of the data for the parameterization in each row. This is useful for model comparison and model criticism but not important for our current adventures.
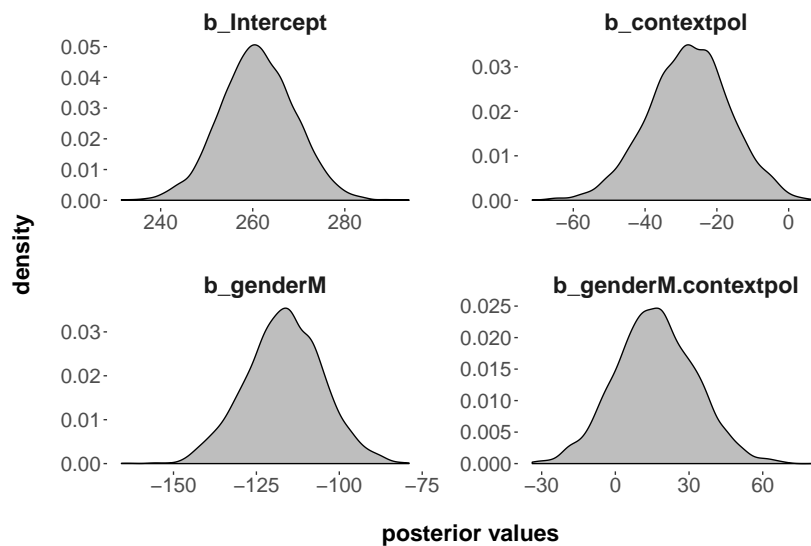

Figure 4: Posterior density of parameter values in the fixed-effects regression model

Now, here comes a nice gadget. Based on the samples obtained for `contextpol` (β_{pol}), it is very easy to estimate our belief that β_{pol} is indeed negative. We simply have to calculate the proportion of samples that were negative, that's all. For instance, with the code below, which reveals that the posterior probability, given the data, that $\beta_{pol} < 0$ is about 0.99275, so very close to 1!

```
# proportion of negative samples for parameter p_contextpol
# this number approximates P(beta_pol < 0 | model, data)
mean(post_samples_FE$b_contextpol < 0)
```

As an interim summary, we have seen how to run a Bayesian regression analysis with the `brms` package and deal with its output. We have also seen

that the output can be interpreted in very intuitive ways (e.g., "The probabil-ity of H1, given our model, priors, and data, is more than .99").

Unfortunately, what we have not seen yet is what our model and data say about hypotheses 2 or 3. This is because there is no single parameter in the (dummy-coded) regression model that corresponds to the differences between cells 3 and 4 (for hypothesis 2) and cells 2 and 3 (for hypothesis 3). Notice that this problem is not specific to Bayesian analyses, but inherent in the way the regression coefficients were set up. However, unlike in more frequentist/classical analysis, the Bayesian approach allows to recover information about any derived measure from the obtained samples. Here's how:

Take hypothesis 3 which requires us to compare cells 2 and 3. The hypoth-esis states that $\beta_0 + \beta_{\mathrm{pol}} > \beta_0 + \beta_{\mathrm{male}}$, which reduces to $\beta_{\mathrm{pol}} > \beta_{\mathrm{male}}$. We can approximate the posterior probability that this is true based on the samples that we obtained for our model in the same general way as before, namely [tr: I think this would be easier to follow if we'd keep the reference level in the equation]:

A potential way of testing different hypothe-ses of the kind we have set our here, is to run different regression analyses, each with a different reference cell. This is a rather unhandy work flow. It wouldn't help with hypothesis 3 either, which compares the cell means in our design matrix "diagonally": there is no way of changing the reference level of either factor such that dummy coding gives us a single coefficient as the difference between cells 2 and 3.

```
1 # proportion of samples where the mean for cell 2 was bigger
2 # than that of cell 3
3 # this number approximates P(beta_pol > beta_male | model, data)
4 mean(post_samples_FE$b_contextpol - post_samples_FE$b_genderM > 0)
```

Based on the posterior samples we obtained, this estimate is 1. That's a strong result. If the model were true, then, given the data, our certainty that hypothesis 3 is true should be pretty much almost at ceiling.

To conclude this section, the Bayesian approach to regression modeling allows us to retrieve all direct comparisons between cells in a factorial design (all possible comparisons). It also allows us to retrieve quantitative infor-mation about our hypotheses which is accessible, and easy to communicate and understand. We can calculate (estimated) posterior probability that a particular hypothesis holds.

The `faintr` *package*

To facilitate the comparison of pairs of cells also in bigger factorial designs, this tutorial comes with a little R package, the `faintr` package. You can install the package from GitHub with the `devtools` package, as follows:

The name `faintr` is indicative of the possibility that the package might break down unexpectedly (we might consider renaming after more extensive testing), but also alludes to "*fa*ctorial design" and "*int*erpretation" somehow.

```
1 # package to allow installation from github
2 library(devtools)
3 # package with convenience function for Bayesian regression
      models for factorial designs
4 install_github("michael-franke/bayes_mixed_regression_tutorial/
      faintr") # install from GitHub
5 library(faintr)
```

The `faintr` package provides two (hopefully) useful functions. The

function `extract_posterior_cell_means` takes as input the output of a factorial-design regression model fitted with brm. It outputs samples for all design cell means, and a comparison of all design cells against each other. The function `get_cell_comparison` takes the same kind of model fit as input, together with a specification of which two cells to compare against each other. Using the latter function, the source code provides a convenient function to produce the posterior probability of the three hypothesis relevant for this tutorial:

```
> get_posterior_beliefs_about_hypotheses_new(modelFE)
# A tibble: 3 x 2
  hypothesis                        probability
  <chr>                                   <dbl>
1 Female-polite < Female-informal         0.993
2 Male-polite < Male-informal             0.842
3 Male-informal < Female-polite           1
```

Adding random effects structures

Although the results look quite different, running hierarchical models with random effects with brms is very similar to the look and feel of non-Bayesian approaches. Here is a function call to a model with the maximal random effect structure licensed by the design:

```
# hierarchical model with the maximial RE structure licensed by
    the design
# (notice that factor 'gender' does not vary for a given value
    of variable 'subject')
model_MaxRE = brm(formula = pitch ~ gender * context +
                    (1 + gender * context | sentence) +
                    (1 + context | subject),
                data = politedata,
                control = list(adapt_delta = 0.9))
```

The outcome of this is model fit is shown here:

```
1   > model_MaxRE
2    Family: gaussian
3     Links: mu = identity; sigma = identity
4   Formula: pitch ~ gender * context + (1 + gender * context | sentence) + (1 + context | subject)
5      Data: politedata (Number of observations: 83)
6   Samples: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
7            total post-warmup samples = 4000
8
9   Group-Level Effects:
10  ~sentence (Number of levels: 7)
11                                  Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
12  sd(Intercept)                      21.84      9.72     7.40    44.90       2116 1.00
13  sd(genderM)                        11.13      9.07     0.49    33.29       2435 1.00
14  sd(contextpol)                     15.04     11.24     0.55    42.49       1637 1.00
15  sd(genderM:contextpol)             16.55     13.43     0.76    47.63       2173 1.00
16  cor(Intercept,genderM)             -0.23      0.44    -0.90     0.71       4808 1.00
17  cor(Intercept,contextpol)           0.01      0.41    -0.74     0.79       4658 1.00
18  cor(genderM,contextpol)            -0.06      0.44    -0.83     0.77       2940 1.00
19  cor(Intercept,genderM:contextpol)  -0.10      0.43    -0.84     0.73       5216 1.00
20  cor(genderM,genderM:contextpol)    -0.03      0.44    -0.82     0.80       3335 1.00
21  cor(contextpol,genderM:contextpol) -0.15      0.44    -0.87     0.74       3057 1.00
22
23  ~subject (Number of levels: 6)
24                       Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
25  sd(Intercept)           36.11     18.30    14.60    84.84       1210 1.00
26  sd(contextpol)           9.17      8.73     0.32    32.53       2290 1.00
27  cor(Intercept,contextpol) 0.03     0.58    -0.93     0.95       4608 1.00
28
29  Population-Level Effects:
30                  Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
31  Intercept         261.92     25.73   213.11   312.17       1088 1.00
32  genderM          -116.78     35.19  -188.21   -50.38       1011 1.00
33  contextpol        -27.16     12.44   -51.30    -2.22       2782 1.00
34  genderM:contextpol  15.13     16.84   -17.62    47.40       2863 1.00
35
36  Family Specific Parameters:
37        Estimate Est.Error l-95% CI u-95% CI Eff.Sample Rhat
38  sigma    24.96      2.38    20.73    30.07       3676 1.00
39
40  Samples were drawn using sampling(NUTS). For each parameter, Eff.Sample
41  is a crude measure of effective sample size, and Rhat is the potential
42  scale reduction factor on split chains (at convergence, Rhat = 1).
```

The lines 29–34 again give the estimates of the fixed-effects coefficients, but we now also obtain information about the parameters implied by the specified random effect structure. Lines 9–21 cover the by-sentence random effects, lines 23–27 cover the by-subject random effects. We see from the 95% credible intervals that the only parameters implied by the random effects structure that does seem to receive sufficient a posteriori credence on non-trivial values are the by-sentence and by-subject random intercepts.

As for the probability of the hypotheses of interest, we can use the faintr package and the convenience function defined above to inspect:

```
> get_posterior_beliefs_about_hypotheses_new(model_MaxRE)
# A tibble: 3 x 2
  hypothesis                    probability
  <chr>                              <dbl>
1 Female-polite < Female-informal    0.981
2 Male-polite < Male-informal        0.830
3 Male-informal < Female-polite      0.988
```

Priors

... in progress ...

References

Buerkner, Paul-Christian (2016). "brms: An R package for Bayesian multi-level models using Stan". In: *Journal of Statistical Software* 80.1, pp. 1–28.

R Core Team (2017). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. URL: https://www.R-project.org/.

Sorensen, Tanner, Sven Hohensteinb, and Shravan Vasishth (2016). "Bayesian linear mixed models using Stan: A tutorial for psychologists, linguists, and cognitive scientists". In: *The Quantitative Methods for Psychology*.

Winter, Bodo (2013). *Linear models and linear mixed effects models in R with linguistic applications*. URL: https://arxiv.org/abs/1308.5499.

Winter, Bodo and S. Grawunder (2012). "The Phonetic Profile of Korean Formality". In: *Journal of Phonetics* 40, pp. 808–815.

Bayesian inference: priors, likelihoods and posteriors

Jones is a rational scientist. She has recently inherited her grandma's lucky coin. Grandma used this coin many times during Jones' childhood to determine whether Jones was allowed a sweet or not. Jones suspects that grandma's coin might be a trick coin, but she is not sure. She is determined to find out. How? Well, naturally, by rationally updating her *prior beliefs* about the coin's bias to obtain a new *posterior belief* based on empirical observation (outcomes of coin flips). Central to this updating is Jones' *likelihood function*, which encodes how likely each relevant coin bias may have generated the observed data. — Sounds fancifully abstract? It's actually fairly intuitive. Consider this example.

Prior beliefs.   Jones initially believes that the coin is either biased towards heads or biased towards tails, and that both of these possibilities are equally likely. She also believes that, if biased towards heads, the coin is two times more likely to come up heads; and that, if biased towards tails, the coin is two times more likely to come up tails. Numerically, Jones' *prior beliefs* can be written as, where $\theta \in [0; 1]$ is the coin's bias: $P(\theta = 1/3) = 1/2$, and $P(\theta = 2/3) = 1/2$.

Likelihood.   The bias θ is, by definition, the probability of the coin landing heads on the next trial. Let's assume that Jones tosses the coin only once (hm, maybe not so rational a scientist after all? or just too busy?). Let D be the set of potential outcomes of this experiment, namely $D = \{\text{heads}, \text{tails}\}$. The *likelihood function* determines the likelihood of observing each observation $d \in D$ for each θ, which in our case is just rather trivial: $P(D = \text{heads} \mid \theta) = \theta$ and $P(D = \text{tails} \mid \theta) = 1 - \theta$.

Posterior beliefs.   Jones observes that the coin landed heads. What should she believe now? By *Bayes rule* her posterior beliefs are defined as:

$$P(\theta \mid D = \text{heads}) = \frac{P(\theta)P(D = \text{heads} \mid \theta)}{\sum_{\theta'} P(\theta')P(D = \text{heads} \mid \theta')}$$

Jones' posterior belief that the coin is twice as likely to land heads is therefore:

$$P(\theta = 2/3 \mid D = \text{heads}) =$$

$$\frac{P(\theta = 2/3)\ P(D = \text{heads} \mid \theta = 2/3)}{P(\theta = 2/3)\ P(D = \text{heads} \mid \theta = 2/3) + P(\theta = 1/3)\ P(D = \text{heads} \mid \theta = 1/3)} =$$

$$\frac{1/2\ 2/3}{1/2\ 2/3 + 1/2\ 1/3} = 1/2$$

After making her observation, rational Jones believes that the bias towards heads is twice more likely than the bias towards tails. [tr: I think what can be confusing it that her prior is a bias of observing heads TWICE as likely as tails (or vice versa) and her posterior is to believe that the one bias is TWICE as likely as the other. If the numbers would be different, we would avoid a possible confusion of prior and posterior. My first reading was: so the posterior didnt change the prior...oh nuh, now I gottit.]

Info Box 1: Priors, likelihood and posteriors in Bayesian inference.

More about the information displayed in the summary output of a brm *model fit*

The first column *Estimate* gives the mean of the obtained samples, thereby approximating the mean of the posterior distribution (beliefs we should hold) about each parameter. For example, the parameter Intercept is estimated to have a mean of about 261, which is (of course) exactly what we calculated as a mean of the data points in cell 1, as shown in Figure 2 [tr: I dont think it is good to give the impression that the mean and the estimate should be identical. Depending on the sampling, the mean might be off the estimate]. The other columns give further useful information. *Est.Error* is the estimation error, an indication of the certainty we should have about the whole inference procedure. The columns l-95% CI and u-95% CI give the lower and upper bound of the **95% credible interval** for each parameter, estimated from the posterior samples.

The column Eff.Sample, for efficient samples, gives a rough measure of how many of all the samples we took (4000 in our case) are contributing non-redundant information to our estimation. The higher this number, the better. Finally, Rhat is a measure of whether the samples obtained are likely representative of the true distribution. Concretely, it indicates whether the four chains we ran in parallel all ended up with the same results, so to speak. If this column contains values bigger than 1.1 this is an indication that your model fit has not converged. (If your model output indicates non-convergence, you may want to increase the number of samples, but you should also consider the possibility that you are trying to fit a model which cannot be "trained" based on the (perhaps insufficient) data and the particular method of posterior sampling. For common regression analyses, this will usually entail considering a simpler model (e.g., with fewer explanatory factors, less (correlated) random effects, etc.))

Info Box 2: Information in summaries of brm model fits.